

Program :

```
package PackageDemo;

import java.io.IOException;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.GenericOptionsParser;

public class ProcessUnits {
    public static void main(String[] args) throws Exception {
        Configuration c = new Configuration();
        String[] files = new GenericOptionsParser(c, args).getRemainingArgs();
        Path input = new Path(files[0]);
        Path output = new Path(files[1]);
        Job j = Job.getInstance(c, "maxconsumption");
        j.setJarByClass(ProcessUnits.class);
        j.setMapperClass(MapForMaxConsumption.class);
        j.setReducerClass(ReduceForMaxConsumption.class);
        j.setOutputKeyClass(Text.class);
        j.setOutputValueClass(Text.class);
        FileInputFormat.addInputPath(j, input);
        FileOutputFormat.setOutputPath(j, output);
        System.exit(j.waitForCompletion(true) ? 0 : 1);
    }

    public static class MapForMaxConsumption extends Mapper<LongWritable, Text, Text, Text> {
        public void map(LongWritable key, Text value, Context con) throws IOException,
            InterruptedException {
            String line = value.toString();
            String[] tokens = line.split(",");
            if (tokens.length != 13) {
                System.err.println("Skipping malformed line: " + line);
                return; // Skip this line
            }
            String year = tokens[0];
            int maxConsumption = Integer.MIN_VALUE;
            String maxMonth = "";
            for (int i = 1; i <= 12; i++) {
                try {
                    int consumption = Integer.parseInt(tokens[i]);
                    if (consumption > maxConsumption) {
                        maxConsumption = consumption;
                    }
                } catch (NumberFormatException e) {
                    // Ignore
                }
            }
        }
    }
}
```

```

        maxMonth = getMonthName(i);
    }
} catch (NumberFormatException e) {
    System.err.println("Skipping invalid consumption value: " + tokens[i] + " for year " +
year);
}
}
}
if (maxMonth.isEmpty()) {
    return; // Skip this entry if no valid max month was found
}
con.write(new Text(year), new Text(maxMonth + ":" + maxConsumption));
}
private String getMonthName(int monthIndex) {
    String[] months = {
        "January", "February", "March", "April", "May", "June",
        "July", "August", "September", "October", "November", "December"
    };
    return months[monthIndex - 1]; // Adjusting for 1-based month index
}
}
public static class ReduceForMaxConsumption extends Reducer<Text, Text, Text, Text> {
    private boolean headerWritten = false; // Flag to check if header is written

    public void reduce(Text year, Iterable<Text> values, Context con) throws IOException,
InterruptedException {
        String maxMonth = "";
        int maxConsumption = Integer.MIN_VALUE;
        for (Text value : values) {
            String[] monthAndConsumption = value.toString().split(":");
            if (monthAndConsumption.length != 2) continue; // Skip if not properly formatted

            String month = monthAndConsumption[0];
            int consumption = Integer.parseInt(monthAndConsumption[1]);
            if (consumption > maxConsumption) {
                maxConsumption = consumption;
                maxMonth = month;
            }
        }
        if (!maxMonth.isEmpty()) {
            if (!headerWritten) {
                // Write the header only once
                con.write(new Text("Year"), new Text("Month"));
                con.write(new Text("Max Consumption"), new Text(""));
                headerWritten = true;
            }
            String outputLine = String.format("Year: %-10s Month: %-15s Max
Consumption: %d", year.toString(), maxMonth, maxConsumption);
            con.write(new Text(""), new Text(outputLine)); // Write formatted output
        }
    }
}
}
}
}

```

Output :

```
[cloudera@quickstart 33316]$ hadoop jar Ass2.jar PackageDemo.ProcessUnits Sample.txt
Assig2Dir1
25/01/21 23:44:05 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
25/01/21 23:44:06 INFO input.FileInputFormat: Total input paths to process : 1
25/01/21 23:44:07 INFO mapreduce.JobSubmitter: number of splits:1
25/01/21 23:44:07 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1737525146805_0003
25/01/21 23:44:07 INFO impl.YarnClientImpl: Submitted application
application_1737525146805_0003
25/01/21 23:44:07 INFO mapreduce.Job: The url to track the job:
http://quickstart.cloudera:8088/proxy/application_1737525146805_0003/
25/01/21 23:44:07 INFO mapreduce.Job: Running job: job_1737525146805_0003
25/01/21 23:44:18 INFO mapreduce.Job: Job job_1737525146805_0003 running in uber mode :
false
25/01/21 23:44:18 INFO mapreduce.Job: map 0% reduce 0%
25/01/21 23:44:26 INFO mapreduce.Job: map 100% reduce 0%
25/01/21 23:44:36 INFO mapreduce.Job: map 100% reduce 100%
25/01/21 23:44:37 INFO mapreduce.Job: Job job_1737525146805_0003 completed successfully
25/01/21 23:44:37 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=91
        FILE: Number of bytes written=220899
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=326
        HDFS: Number of bytes written=333
        HDFS: Number of read operations=6
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=1
        Launched reduce tasks=1
        Data-local map tasks=1
        Total time spent by all maps in occupied slots (ms)=6624
        Total time spent by all reduces in occupied slots (ms)=7778
        Total time spent by all map tasks (ms)=6624
        Total time spent by all reduce tasks (ms)=7778
        Total vcore-seconds taken by all map tasks=6624
        Total vcore-seconds taken by all reduce tasks=7778
        Total megabyte-seconds taken by all map tasks=6782976
        Total megabyte-seconds taken by all reduce tasks=7964672
    Map-Reduce Framework
        Map input records=5
        Map output records=5
        Map output bytes=75
        Map output materialized bytes=91
        Input split bytes=121
        Combine input records=0
        Combine output records=0
```

Reduce input groups=5
Reduce shuffle bytes=91
Reduce input records=5
Reduce output records=7
Spilled Records=10
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=135
CPU time spent (ms)=1200
Physical memory (bytes) snapshot=357376000
Virtual memory (bytes) snapshot=3007225856
Total committed heap usage (bytes)=226365440

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=205

File Output Format Counters

Bytes Written=333

[cloudera@quickstart 33316]\$ hadoop fs -cat Assig2Dir1/part-r-000000

Year: 1979	Month: September	Max Consumption: 68
Year: 1980	Month: August	Max Consumption: 35
Year: 1981	Month: May	Max Consumption: 39
Year: 1984	Month: December	Max Consumption: 48
Year: 1985	Month: June	